

Team name: greedybluemoon

Group member: Yichen Zhou, Feifan Zhang

netid: zhou113, fzhang19

Milestone 1:

list of all kernels that collectively consume more than 90% of the program time:

void fermiPlusCgemmLDS128_batched: 34.18%

void cudnn::detail::implicit_convolve_sgemm: 27.09%

void fft2d_c2r_32x32: 12.71%

Sgemm_sm35_ldg_tn_128x8x256x16x32: 8.25%

[CUDA memcpy HtoD]: 6.49%

void cudnn::detail::activation_fw_4d_kernel: 4.08%

list of all CUDA API calls that collectively consume more than 90% of the program time.

cudaStreamCreateWithFlags: 37.37%

cudaFree: 28.82%

cudaMemGetInfo: 23.62%

cudaStreamSynchronize: 8.37%

Include an explanation of the difference between kernels and API calls

A kernel is a low level program interfacing with hardware. It is the lowest level program running

on computers. An API is a generic term defining the interface developers have to use when writing code using libraries and a programming language. So when we are using API, it is the same as we call the code which are being wrote by other programmer.

Show output of rai running MXNet on the CPU

```
Loading fashion-mnist data...
done
Loading model...
done
New Inference
EvalMetric: {'accuracy': 0.8444}
12.62user 6.40system 0:08.46elapsed 224%CPU (0avgtext+0avgdata 2828716maxresident)k
0inputs+2624outputs (0major+39582minor)pagefaults 0swaps
```

List program run time

```
12.62user 6.40system 0:08.46elapsed 224%CPU (0avgtext+0avgdata 2828716maxresident)k
0inputs+2624outputs (0major+39582minor)pagefaults 0swaps
```

Show output of rai running MXNet on the GPU

```

Loading fashion-mnist data...
done
Loading model...
==342== NVPROF is profiling process 342, command: python m1.2.py
[18:17:30] src/operator/././cudnn_algoreg-inl.h:112: Running performance tests to find
the best convolution algorithm, this can take a while... (setting env variable
MXNET_CUDNN_AUTOTUNE_DEFAULT to 0 to disable)
done
New Inference
EvalMetric: {'accuracy': 0.8444}
==342== Profiling application: python m1.2.py
==342== Profiling result:
Time(%)    Time    Calls    Avg      Min      Max  Name
33.90%  118.53ms      9  13.170ms  13.158ms  13.188ms  void
fermiPlusCgemmLDS128_batched<bool=0, bool=1, bool=0, bool=0, int=4, int=4,
int=3, int=3, bool=1, bool=1>(float2**, float2**, float2**, float2*, float2 const *,
float2 const *, int, int, int, int, int, __int64, __int64, __int64, float2 const
*, float2 const *, float2, float2, int)
26.87%  93.944ms      1  93.944ms  93.944ms  93.944ms  void
cudnn::detail::implicit_convolve_sgemm<float, int=1024, int=5, int=5, int=3, int=3,
int=3, int=1, bool=1, bool=0, bool=1>(int, int, int, float const *, int,
cudnn::detail::implicit_convolve_sgemm<float, int=1024, int=5, int=5, int=3, int=3,
int=3, int=1, bool=1, bool=0, bool=1>*, float const *, kernel_conv_params, int, float,
float, int, float const *, float const *, int, int)
12.63%  44.153ms      9  4.9059ms  2.7041ms  6.2815ms  void fft2d_c2r_32x32<float,
bool=0, unsigned int=0, bool=0, bool=0>(float*, float2 const *, int, int, int, int,
int, int, int, int, int, float, float, cudnn::reduced_divisor, bool, float*, float*)
8.16%  28.540ms      1  28.540ms  28.540ms  28.540ms
sgemm_sm35_ldg_tn_128x8x256x16x32
6.87%  24.035ms     14  1.7168ms  1.5360us  23.208ms  [CUDA memcpy HtoD]
4.06%  14.178ms      2  7.0892ms  252.57us  13.926ms  void
cudnn::detail::activation_fw_4d_kernel<float, float, int=128, int=1, int=4,
cudnn::detail::tanh_func<float>>(cudnnTensorStruct, float const *,

```

```

cudnn::detail::activation_fw_4d_kernel<float, float, int=128, int=1, int=4,
cudnn::detail::tanh_func<float>>, cudnnTensorStruct*, float, cudnnTensorStruct*, int,
cudnnTensorStruct*)
  3.80% 13.297ms      1 13.297ms 13.297ms 13.297ms void
cudnn::detail::pooling_fw_4d_kernel<float, float,
cudnn::detail::maxpooling_func<float, cudnnNanPropagation_t=0>,
int=0>(cudnnTensorStruct, float const *, cudnn::detail::pooling_fw_4d_kernel<float,
float, cudnn::detail::maxpooling_func<float, cudnnNanPropagation_t=0>, int=0>,
cudnnTensorStruct*, cudnnPoolingStruct, float, cudnnPoolingStruct, int,
cudnn::reduced_divisor, float)
  1.71% 5.9897ms      9 665.52us 499.99us 886.25us void fft2d_r2c_32x32<float,
unsigned int=0, bool=0>(float2*, float const *, int, int, int, int, int, int, int,
int, int, cudnn::reduced_divisor, bool)
  1.16% 4.0654ms      1 4.0654ms 4.0654ms 4.0654ms
sgemm_sm35_ldg_tn_64x16x128x8x32
  0.37% 1.2819ms      1 1.2819ms 1.2819ms 1.2819ms void
mshadow::cuda::MapPlanLargeKernel<mshadow::sv::saveto, int=8, int=1024,
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>,
mshadow::expr::Plan<mshadow::expr::ScalarExp<float>, float>>(mshadow::gpu, unsigned
int, mshadow::Shape<int=2>, int=2, int)
  0.32% 1.1108ms      1 1.1108ms 1.1108ms 1.1108ms void
mshadow::cuda::SoftmaxKernel<int=8, float,
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>,
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>>(mshadow::gpu,
int=2, unsigned int)
  0.05% 175.84us      13 13.525us 2.0480us 74.590us void
mshadow::cuda::MapPlanKernel<mshadow::sv::saveto, int=8,
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>,
mshadow::expr::Plan<mshadow::expr::ScalarExp<float>, float>>(mshadow::gpu, unsigned
int, mshadow::Shape<int=2>, int=2)
  0.04% 146.43us      2 73.214us 16.256us 130.17us void
mshadow::cuda::MapPlanKernel<mshadow::sv::plusto, int=8,
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>,
mshadow::expr::Plan<mshadow::expr::Broadcast1DExp<mshadow::Tensor<mshadow::gpu, int=1,
float>, float, int=2, int=1>, float>>(mshadow::gpu, unsigned int,
mshadow::Shape<int=2>, int=2)
  0.04% 130.46us      1 130.46us 130.46us 130.46us
sgemm_sm35_ldg_tn_32x16x64x8x16
  0.01% 22.495us      1 22.495us 22.495us 22.495us void
mshadow::cuda::MapPlanKernel<mshadow::sv::saveto, int=8,
mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=2, float>, float>,
mshadow::expr::Plan<mshadow::expr::ReduceWithAxisExp<mshadow::red::maximum,
mshadow::Tensor<mshadow::gpu, int=3, float>, float, int=3, bool=1, int=2>,
float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=2)
  0.01% 21.344us      1 21.344us 21.344us 21.344us void fft2d_r2c_32x32<float,
unsigned int=5, bool=1>(float2*, float const *, int, int, int, int, int, int, int,

```

```

int, int, cudnn::reduced_divisor, bool)
0.00% 9.9510us 1 9.9510us 9.9510us 9.9510us [CUDA memcpy DtoH]
==342== API calls:
Time(%) Time Calls Avg Min Max Name
37.40% 1.45873s 18 81.041ms 22.761us 729.21ms cudaStreamCreateWithFlags
29.25% 1.14096s 10 114.10ms 1.2160us 302.81ms cudaFree
23.47% 915.65ms 27 33.913ms 259.70us 907.34ms cudaMemGetInfo
8.30% 323.71ms 29 11.162ms 6.5270us 194.11ms cudaStreamSynchronize

1.24% 48.507ms 9 5.3897ms 8.6400us 23.394ms cudaMemcpy2DAsync
0.18% 7.0216ms 45 156.04us 9.0380us 966.62us cudaMalloc
0.04% 1.3659ms 4 341.46us 336.29us 356.33us cuDeviceTotalMem
0.03% 1.2067ms 352 3.4280us 516ns 113.28us cuDeviceGetAttribute
0.03% 1.0783ms 114 9.4580us 1.2580us 406.41us cudaEventCreateWithFlags
0.02% 788.60us 53 14.879us 6.4670us 95.007us cudaLaunch
0.01% 364.98us 6 60.829us 27.318us 119.28us cudaMemcpy
0.01% 356.02us 619 575ns 524ns 1.5840us cudaSetupArgument
0.01% 289.79us 4 72.447us 52.258us 92.095us cudaStreamCreate
0.00% 124.11us 116 1.0690us 680ns 2.5100us cudaDeviceGetAttribute
0.00% 104.53us 4 26.133us 19.246us 29.666us cuDeviceGetName
0.00% 93.609us 36 2.6000us 826ns 8.4420us cudaSetDevice
0.00% 56.850us 27 2.1050us 1.7510us 3.6070us cudaStreamWaitEvent
0.00% 51.082us 53 963ns 533ns 2.7090us cudaConfigureCall
0.00% 47.488us 2 23.744us 23.146us 24.342us
cudaStreamCreateWithPriority
0.00% 26.169us 10 2.6160us 1.5350us 8.3280us cudaGetDevice
0.00% 23.543us 12 1.9610us 1.1230us 5.2820us cudaEventRecord
0.00% 23.013us 1 23.013us 23.013us 23.013us cudaBindTexture
0.00% 21.063us 34 619ns 537ns 1.0390us cudaGetLastError
0.00% 13.161us 18 731ns 638ns 1.0900us cudaPeekAtLastError
0.00% 7.1140us 6 1.1850us 572ns 2.3200us cuDeviceGetCount
0.00% 6.5660us 6 1.0940us 645ns 1.6480us cuDeviceGet
0.00% 6.5200us 1 6.5200us 6.5200us 6.5200us cudaStreamGetPriority
0.00% 6.3120us 1 6.3120us 6.3120us 6.3120us cudaEventCreate
0.00% 5.0090us 2 2.5040us 2.0170us 2.9920us
cudaDeviceGetStreamPriorityRange
0.00% 3.2640us 3 1.0880us 972ns 1.1480us cuInit
0.00% 2.8040us 3 934ns 889ns 1.0150us cuDriverGetVersion
0.00% 2.5250us 1 2.5250us 2.5250us 2.5250us cudaEventDestroy
0.00% 1.8790us 1 1.8790us 1.8790us 1.8790us cudaUnbindTexture
0.00% 1.3280us 1 1.3280us 1.3280us 1.3280us cudaGetDeviceCount

```

List program run time

```
2.31user 1.07system 0:02.84
```

Milestone 2:

```
Loading fashion-mnist data...
done
Loading model...
done
New Inference
Op Time: 6.585881
Op Time: 19.545969
Correctness: 0.8451 Model: ece408
30.66user 1.58system 0:30.12elapsed 107%CPU (0avgtext+0avgdata 2819428maxresiden
t)k
```

Total Execution Time:

$30.66 + 1.58 + 30.12 = 62.36s$

Operation Times:

First layer Op Time: 6.585881s

Second layer Op Time: 19.545969s

Milestone 3:

In our parallel implemented forward_kernel, its duration is about 705 microseconds, while in the sequential kernel in milestone 2 it costs about 20 seconds. And at the same time, they have the exact same accuracy and same functionality.

mxnet::op::forward_kernel(float*, float const *, float const *, int, int, int, int, int, int)

Duration	704.981 μ s
----------	-----------------

We didn't use shared memory in our forward kernel, so the optimization effect of shared memory doesn't appear in our case.

Shared Memory/Block	0 B
---------------------	-----

The following is the data for the kernel function we used:

Cuda Functions :

<code>mxnet::op::forward_kernel(float*, float const *, float const *, int, int, int, int, int, int)</code>
--

Maximum instruction execution count in assembly: 249600

Average instruction execution count in assembly: 72873

Instructions executed for the kernel: 9182080

Thread instructions executed for the kernel: 239962880

Non-predicated thread instructions executed for the kernel: 224787840

Warp non-predicated execution efficiency of the kernel: 76.5%


Warp execution efficiency of the kernel: 81.7%

From this we observed that the kernel's warp execution efficiency of 76.5% is less than 100% due to divergent branches and predicated instructions. If predicated instructions are not taken into account the warp execution efficiency for these kernels is 81.7%. This could lead to lower utilization of the GPU's compute resources. The way to optimize this is through reducing the amount of intra-warp divergence and predication in the kernel.

Another thing is the divergent branches. If there exists a high divergence in some warps, the total efficiency of execution would decrease dramatically since divergence means waste of execution resource in each warps. So if we could rearrange each warps and threads to make consecutive threads executing or not executing at the same time, the divergence would be reduced which means that higher optimization in parallel computing kernels.

Memory bandwidth is another important factor that restricts the performances of parallel kernel. But since we didn't use shared memory here, we didn't have the best optimized performance due to the restriction in speed of global memory reads and

global memory writes.


Transactions	Bandwidth	Utilization	
L1/Shared Memory			
Local Loads	0	0 B/s	
Local Stores	0	0 B/s	
Shared Loads	0	0 B/s	
Shared Stores	0	0 B/s	
Global Loads	2062544	166.377 GB/s	
Global Stores	8840	885.478 MB/s	
Atomic	0	0 B/s	
L1/Shared Total	2071384	167.262 GB/s	


Milestone 4:

Our first optimization is the Unroll / shared-memory Matrix multiply. In this function we have two kernels. The first is the unroll function to expand the input feature map into a unrolled version. And then the next step is to use the matrix multiplication function to multiply the X_unrolled with the w input. Shared memory is used in the matrix multiplication.

Also when running dataset of 10, it showed worse performance than the restricted and normal ones.

Here is the timeline of the unroll method.





I first examined the matrix multiplication kernel.

Duration	31.776 μ s
Grid Size	[43,1,1]
Block Size	[16,16,1]
Registers/Thread	25
Shared Memory/Block	2 KiB
Shared Memory Requested	112 KiB
Shared Memory Executed	112 KiB
Shared Memory Bank Size	4 B

According to the analysis, our grid size is too small to hide compute and memory latency. The kernel does not execute enough blocks to hide memory and operation latency. Typically the kernel grid size must be large enough to fill the GPU with multiple "waves" of blocks. Since the grid size is determined by the height and width of output and number of maps. This could be much more efficient when dealing with big datasets.

With the use of shared memory, we do not need to perform global read for this kernel. The shared loads and stores are used efficiently.

Local Loads	0	0 B/s	
Local Stores	0	0 B/s	
Shared Loads	82560	577.769 GB/s	
Shared Stores	6880	48.147 GB/s	
Global Loads	18477	29.707 GB/s	
Global Stores	940	1.484 GB/s	
Atomic	0	0 B/s	
L1/Shared Total	108857	657.107 GB/s	<div> <div></div> <div>Idle</div> <div>Low</div> <div>Medium</div> <div>High</div> <div>Max</div> </div>

The next is the unroll kernel. This changed the X input matrix into an unroll matrix. Here is the major parameters.

Duration	19.488 μ s
Grid Size	[4,1,1]
Block Size	[1024,1,1]
Registers/Thread	23
Shared Memory/Block	0 B
Shared Memory Requested	112 KiB
Shared Memory Executed	112 KiB
Shared Memory Bank Size	4 B

Our other optimization is using the constant memory and restricted. We use `cudaMemcpyToSymbol` to copy W matrix into a constant memory. The total execution time is 2s this time.



Here is the spec for optimized forward kernel.

Duration	359.484 μ s
Grid Size	[10,16,4]
Block Size	[16,16,1]
Registers/Thread	30
Shared Memory/Block	0 B
Shared Memory Requested	112 KiB
Shared Memory Executed	112 KiB

Since we used constant memory in this scenario, the kernel only performs global stores. So it saved time for global reads.

Local Loads	0	0 B/s	
Local Stores	0	0 B/s	
Shared Loads	0	0 B/s	
Shared Stores	0	0 B/s	
Global Loads	0	0 B/s	
Global Stores	8840	1.744 GB/s	
Atomic	0	0 B/s	
L1/Shared Total	8840	1.744 GB/s	

The third optimization we used is Tuning with restrict, loop unrolling. In fact, we combine this optimization with the constant memory optimization together. First we add the `__restricted__` signal before three input pointers of matrix. So that the compiler would know that writes and reads to these three pointers would not affect the result of another pointers' value. So that the kernel could be accelerated. Also, we use `#pragma unroll` just before the calculation for loop in the `forward_layer` kernel. Thus the compile would unroll all the threads to different threads which would accelerate the operation time. We could know from the picture that the op time is shortened compared to the original

non-optimized kernel.

```
Loading fashion-mnist data...
done
Loading model...
==377== NVPROF is profiling process 377, command: python m4.1.py
done
New Inference
C = 1
B = 10
M = 6
H = 64
K = 5
Op Time: 0.000223
C = 6
B = 10
M = 16
H = 30
K = 5
Op Time: 0.000440
Correctness: 1.0 Model: ece408
==377== Generated result file: /build/timeline.nvprof
* Running nvprof --analysis-metrics -o analysis.nvprof python m4.1.py
```

The two following figures are the screen shot from the NVVP, which is the timeline of running the program by this optimized kernel.

